

# Illuminating a small network: Perspective from within

---

**StrangeBit**

2026



# Contents

<b>Contents</b>	<b>3</b>
<b>1. Introduction</b>	<b>5</b>
1.1 Questions . . . . .	5
<b>2. Background</b>	<b>7</b>
2.1 Basic network characteristics . . . . .	7
2.2 Topologies . . . . .	9
2.3 Tools . . . . .	11
2.4 Network protocols to watch out for . . . . .	12
2.4.1 Spanning trees . . . . .	13
2.4.2 TCP, UDP and ICMP . . . . .	13
2.4.3 DNS and DHCP . . . . .	13
2.4.4 LDAP and other Windows services . . . . .	13
2.4.5 What about timing: NTP . . . . .	13
2.4.6 SSH, TLS and other security protocols . . . . .	13
2.4.7 Watch out for anti-virus . . . . .	13
<b>3. Results</b>	<b>15</b>
3.1 Data collection . . . . .	15
3.2 Breakdown of the network protocols in small enterprise net- work . . . . .	17
3.3 Watch out! We are plotting the network map . . . . .	21
3.4 Performance, performance and once again performance: Stress- ing the outside world . . . . .	21
3.5 Looking for the world: CGNs and TCP . . . . .	21
<b>4. Conclusions</b>	<b>23</b>
<b>Bibliography</b>	<b>25</b>



# 1. Introduction

## 1.1 Questions

In this section we enumerate research questions that we would like to answer in this report. We foresee at least three important questions in this report. Namely those are: **(i) What to measure in the enterprise networks?** By answering this question we attempt to shed the light on most important characteristics in the network traffic analysis. Be it latency, throughput, goodput, error and loss rate, availability. **(ii) How to measure?** Here we would like to answer how to perform the network measurements. For example, how to select vantage point, how to minimize the dataset, but still be able to grasp the most important characteristics of the network. **(iii) When to measure?** This question is also important for a number of reasons. Selecting correct measurement period and correct duration of the measurements intervals will have the important impact on the quality of the research outcome.

By answering these questions properly one can illuminate the performance of the networking infrastructure. Note, in this work we are considering only small networks, comprising 10-100 devices. However, we believe that these questions are also applicable to larger networks and more complicated topologies.



## 2. Background

This section consists of several important parts: First, we discuss main network characteristics of the network; Second, we describe various types of network topologies and device orchestrations; Third, we present various tools which are useful in enterprise network measurements; Finally, we discuss most widely spread network protocols to watch out for in the enterprise network traffic.

### 2.1 Basic network characteristics

We believe that there are several key network characteristics: delay, jitter, throughput, goodput, error and loss rates. All these metrics can be used to measure the performance of the networked systems. In the paragraphs that follow we will describe these metrics and try to explain why they are so important.

*Delay* is the time it takes for the packet (on network layer) to reach the other communication side. Delay can be one way or two way, also called *round trip time*. The former one is hard to measure since the clocks on both sides need to be synchronized. Technically, of course we can use sophisticated algorithms and packet trains to measure one way delay (for example, the reader can look at the RFC 7679 []). But most of the time people rely on half of the RTT. This metric is less accurate since the packets can travel different paths and, hence, delays can be different. But yet this metric is quite common. For example, *ping* utility reports RTT as the measure of delay. We should note that delay impacts user experience greatly, and therefore, it is good to have links with low delays.

*Jitter* is yet another important metric and is widely spread across network engineers. Jitter is the variation of the delay, that is jitter shows how much the delay is varying throughout time. This metric is important

because it can affect how the protocol timers are calculated. For example, if the jitter is high, the calculated timers can be inaccurate and, hence, the performance of such protocols can be undermined. Thus, the lower the jitter the better, in authors opinion, the performance of the networked devices. One easy way to compute jitter is to build the histogram of the network delays and compute the variance.

*Throughput* is also important and it captures how much data (including protocol headers and user payload) can be delivered throughout network system in predefined time interval. Typically, the throughput is measured in Kb/s, Mb/s and Gb/s. Obviously, the larger is the throughput the better network operates. Networks with large throughput can service larger number of clients. *Goodput* is the same as throughput, but excludes the control data from the calculations. In other words, packet header is excluded from the calculations and only user's payload is considered.


*Error rate* describes how often the packets arrive at the receiver with the corrupted bits. Most of the protocols use notion of reliability (that is if the data is corrupted it is requested again) and, hence, high error rate can reduce the performance of such protocols considerably. It is therefore important for the network engineers to avoid highly unstable links. Different media and operational environments have different error rates. For example, wireless links are often have grater error rate than wired and optical links. Also, different applications have different tolerance to errors. For example, Voice over IP and Video over IP require error rates to be low. Mail systems, on the other side, can tolerate high error rates, because the system works in the background and corrupted packets can be requested again.

*Loss rate* is the final metrics that we will cover. Loss occurs, for example, when intermediary routers drop the packet because of congestion and corruption of the packet. Once again, similar to error rates, sensitive applications do not operate well in lossy environments. Hence, typically, network engineers design systems so that such applications will use links with little loss, while other traffic such as HTTP and SNMP protocols can use less expensive, but yet, lossy links. Typically, congested and wireless links with weak signals and obstructed with concrete, for example, have higher loss rates than fat wired and optical links.

There are also other characteristics that can be measured, for example, the reader can take a look at the congestion and some other. But we are not going to cover those in this report.

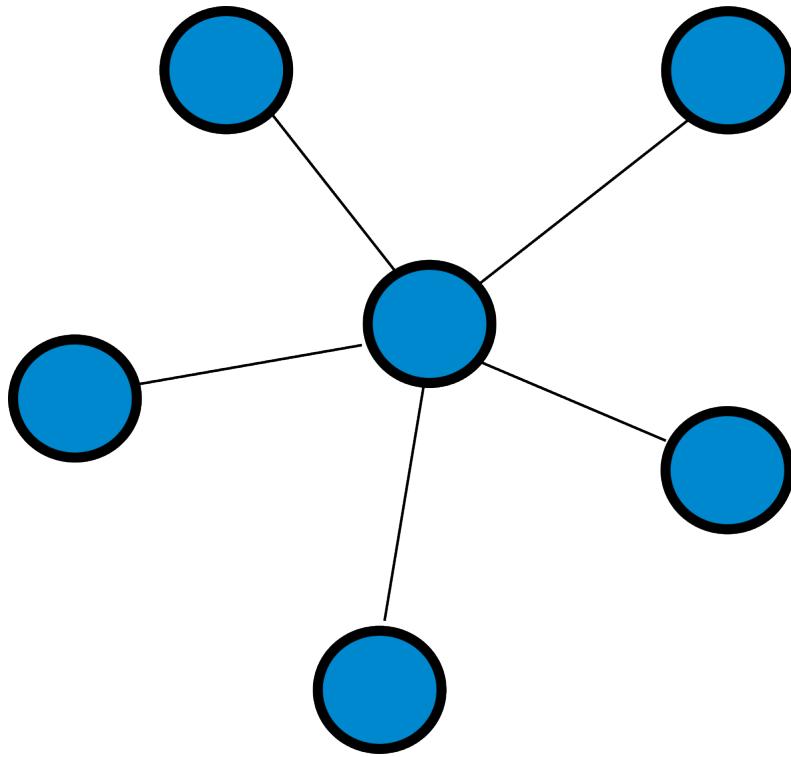
## 2.2 Topologies

There are several key topologies that are used in enterprises: mesh, star and hybrid. *Mesh* topology is such topology in which every network element is connected to every other network element. The links can be wired, wireless and pseudo links (if we are talking about overlays). Mesh can be full and partial. In full mesh, obviously, every element is connected every other element in the network. Consider, for example, personal area network in which all nodes are connected using wireless medium. Wired meshes are expensive, though, and are rarely used in modern deployments. Meshes are crucial, however, when availability is a must. In mesh, some nodes can fail, yet, the network will remain alive and packets will be delivered, not to all, but some devices at least. A typical mesh network topology is shown in Figure 2.1.



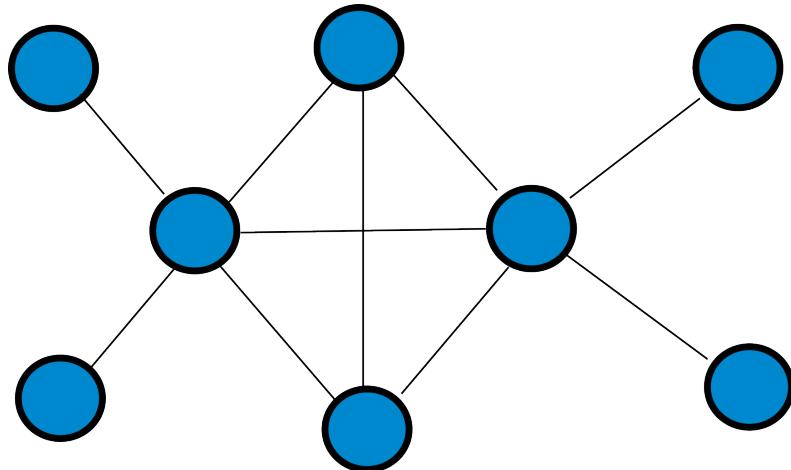
**Figure 2.1.** Mesh topology

*Star* topologies are cheaper and less fault tolerant. In star-like topology some node becomes the root, while others connected to the root element and all the traffic flows through it. Oftentimes, redundant links are added to the topology to bring some level of tolerance to failures. A typical star topology is shown in Figure 2.2.



**Figure 2.2.** Star topology

Finally, there are also hybrid topologies, such as *hub-and-spoke* networks. In this type of networks spoke nodes are connected to hubs, while hubs form a full mesh between each other. Such networks are cheaper than full mesh, but more fault tolerant than star topologies. A typical network is shown in Figure 2.3.



**Figure 2.3.** Hub-and-spoke topology

## 2.3 Tools

Network engineers use wide variety of tools for measuring the performance of the networks and in debugging tasks. All tools can be categorized based on tasks they are meant to be used for such as measuring and troubleshooting.

*Measurement tasks:* to measure performance of the network a common set of tools includes ping utility, tcpping utility, traceroute utility, nc tool, speedtest and iperf, snmp statistics reports by agents and direct queries. In the following paragraph we will describe these tools.

*Ping* utility is the most common tool to measure reachability of a host and measure the round trip times. The tool is based on the ICMP protocol which we will describe later. This tool is widely available in Linux, Windows, BSD and Unix operating systems.

*TCPing* is another common network tool to measure the reachability of the TCP port in the network. But the tool can be also used to measure the time to establish a TCP connection. To our best knowledge this tool is widely available for Windows OS without any charges.

*speedtest* and *iperf* utilities were primarily designed for measuring the bandwidth between two systems. The later one has client and server implementations. This means that to measure the bandwidth (both UDP and TCP connections can be used) one needs to start first the server with the -s flag, and only then run the client with -c flag.

*traceroute* is used often to trace the path the packet takes from host A to host B. The tool uses ICMP under the hood. Not only it is used to trace the reachability of the intermediate routers but it also reports the RTTs.

*netcat*, or *nc* is the tool that is available in Linux distributions and commonly used to send the commands to the server over UDP sockets. It can be used to measure the bandwidth. For example, network administrators can send large enough binary package to the server and measure the time it takes for the transmission to happen.

And finally *SNMP* can be used to collect reports from agents about network performance. Such characteristics as bytes per second delivered by the network interface, ping RTT and many more can all be collected from network devices either with direct queries or with the help of agents.

*Trouble shooting.* A set of tools available in this category. nmap, tcpdump, Wireshark, nslookup, dig, iproute, and ifconfig. All these tools are indispensable in analyzing and troubleshooting the network problems. It

is essential for any network engineer to be acquainted and actively apply these tools. In the paragraphs that follows we will describe what every tool means, but briefly.

*nmap* is a tool that allows network engineer to detect open network ports (both UDP and TCP). Essentially, the tool scans the remote system and reports which port is open. Network engineers and hackers actively use this tool to detect weak points in the network systems.

*tcpdump and Wireshark*. These tools are helpful in collecting and analyzing network protocols. *tcpdump* is typically used to collect the raw packets and frames from the network interface of interest, while *Wireshark* is more advanced it can analyze the network traffic and export it to XML and JSON file formats. We ourselves use these tools in our network traffic analysis tasks.

*nslookup and dig* are used to detect the problems with the DNS servers and queries. These tools are often used by network engineers to detect the problems with DNS and get information about remote systems.

*iproute* is a tool that can be used for multiple purposes, but primarily this tools is useful in configuring the routes to remote systems. For example, network engineers can use the tool to check whether the routes to remote hosts exist. It can be also used to add static routes. But these are just few examples.

And finally, *ifconfig* tool can be used to configure the network interfaces. For example, network engineers can use this tool to set default gateway, IP address on the interface, and set the DNS server name, either manually or with help of DHCP server.

## 2.4 Network protocols to watch out for

There is a humongous amount of protocols in modern networks: from L2 STP to Voice of IP at L4. But we are interested in only few protocols that are common during basic network analysis and troubleshooting: We will cover L2 STP protocol, TCP, UDP, ICMP and some security protocols at L4 such as SSH and TLS.

**2.4.1 Spanning trees**

**2.4.2 TCP, UDP and ICMP**

**2.4.3 DNS and DHCP**

**2.4.4 LDAP and other Windows services**

**2.4.5 What about timing: NTP**

**2.4.6 SSH, TLS and other security protocols**

**2.4.7 Watch out for anti-virus**



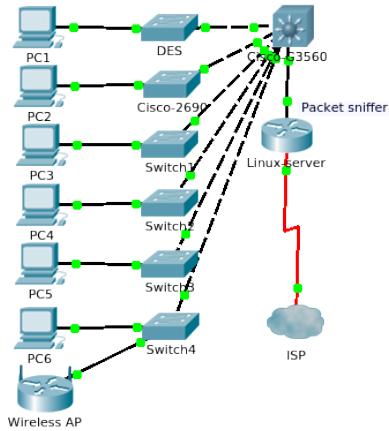
### 3. Results

Analyzing the traffic in the company on regular basis can be beneficial for several reasons: (i) we can find strange behaviour of computers in the network, (ii) we can find the bottlenecks, (iii) we can detect misconfigurations in networked devices. In the paragraphs that follow we present the analysis traffic which was captured in a small enterprise.

#### 3.1 Data collection

We begin this section with a description of the enterprise network, with the details of data capturing process to follow. According to Cisco documentation [?], our network was two-tier medium-sized enterprise network comprising slightly more than 200 workstations, one distribution switch (Cisco Catalyst G3560 series layer 3 switch), one router (the router was running Linux operating system) playing the role of a *VLAN router*, 6 access switches from various vendors (1 DLink DES3200, 1 Cisco 2690, 2 DLink DES1226 and 2 XNET SH9024 switches) and also multiple, non-managed switches and one wireless access point all being connected to access switches.

Since we were unable to analyze fully the wiring of the access (especially those that were spread around the building and connected directly to access switches) and distribution switches, we could basically only guess (based on the limited information we had at our disposal) that there were no redundant links between access and distribution switches, also there were no redundant links between distribution switch and VLAN router. Essentially, the switches were merely connected in a tree like topology: each access switch was connected to one distribution switch, and the distribution switch was connected to a single router over a gigabit Ethernet trunk link.



**Figure 3.1.** Enterprise network architecture

The entire network was partitioned into 25 broadcast domains (VLANs), with native VLAN (VLAN that does not have tagging) being reserved for infrastructure management, although not all the switches were using native VLAN for management purposes. For brevity and privacy reasons we omit the discussion of network addressing in full detail, but rather mention that all VLANs were assigned addresses from class C private network from 192.168.0.0 to 192.168.24.0, whereas the last usable address in each subnetwork was used as the address for the default gateway. All hosts in the network used static address assignment schema.

To capture the traffic it was natural to place packet capturing at the central router, rather than installing it on multiple switches. With this respect, we configured sniffer at the router to capture the frames on a gigabit trunk interface. We have used *tcpdump* to capture all traffic on the trunk link of a router - the link which was connected to the distribution switch. The data collection period lasted for slightly more than 2.5 hours, from 06 : 23 UTC to 09 : 03 UTC and  $19.1 \cdot 10^6$  frames were collected in total. Of course, this is not enough for the full picture (in practice the traffic needs to be collected for at least one week, but this interval was sufficient to understand what kind of traffic existed in the network). Once the data was collected, a single *pcap* file was created. We then started to preprocess the captured data.

**Table 3.1.** Distribution of non IPv4 frames

<b>Protocol</b>	<b>Number of frames</b>	<b>Fraction (%)</b>
ARP	136857	48.8
Cisco Shared Spanning Tree Protocol	119385	42.5
IPv6	17034	6.1
Spanning Tree Protocol (IEEE 802.1D)	4774	1.7
IPX	1164	0.4
Cisco Loop	957	0.3
Cisco CDP/VTP	510	0.2

### 3.2 Breakdown of the network protocols in small enterprise network

To perform the preprocessing we have extracted the binary data and stored it in *pdml* (Packet Description Markup Language) format - a special XML format which is used to describe the contents of the binary representation of frames and packets.

Our next step was to extract needed fields from the XML file and store the extracted data in the format which was simpler for further analysis. Thus, we have chosen JSON (JavaScript Object Notation) format. From PDML we have extracted the following fields:

- Ethernet source and destination MAC addresses and type field
- VLAN identifier and type field
- From IP header we have extracted source and destination addresses, version number, length, protocol type, and TTL.
- Basic information from GRE and PPP headers (since clients accessing the Internet were using these protocols)
- From TCP header we have selected source and destination ports, stream identifier, length, sequence and acknowledgment numbers, as well as flags and window size.
- From UDP header source and destination ports, as well as datagram length information were extracted.

Our next step in preprocessing the data was to extract non IPv4 frames - frames for which neither Ethernet type nor VLAN type field were equal to **0x00008000** (this dataset includes packets from all VLANs as well as packets from native VLAN). Once the data was extracted we binned the frames according to protocol types. In Table 3.1 we present the summary results for this data.

**Table 3.2.** Distribution of protocols in native VLAN

Protocol	Number of frames	Fraction (%)
Cisco Shared Spanning Tree Protocol	4774	40.5
Spanning Tree Protocol (IEEE 802.1D)	4774	40.5
Cisco Loop protocol	957	8.1
IP (ICMP only)	640	5.4
Cisco CDP/VTP	510	4.3
ARP	128	1.1

**Table 3.3.** Frames with invalid MAC addresses

Source MAC	Destination MAC	Source IP	Destination IP	Manufacturer
30:f9:ed:41:a6:01	00:c0:ee:9a:5a:85	192.168.5.10	192.168.5.151	Sony/KYOCERA
00:30:05:c2:b7:ff	00:17:c8:03:a0:7b	192.168.18.7	192.168.18.34	Fujitsu/KYOCERA
00:15:58:67:5f:14	00:c0:ee:9a:5a:85	192.168.5.3	192.168.5.151	FOXCONN/KYOCERA

The next step in preprocessing of the data was to exclude the frames for the native VLAN from the traces. Thus, we have filtered out the frames for which Ethernet type was not equal to **0x00008100**. It turned out that the fraction of frames without VLAN tag was rather small and constituted only 0.06% of the total number of frames in the trace. Since native VLAN was used only for management purposes we excluded it from further analysis. However, in Table 3.2 we show the distribution of packet types seen in the native VLAN.

Instead, next we turned our attention to frames with invalid MAC addresses. By invalid MAC address we mean those MAC addresses which are not multicast or broadcast addresses received at the trunk interface of the router with destination unicast MAC addresses not of router's own MAC addresses. Upon filtering the data we have found that there were two such destination unicast MAC addresses. We hypothesize that such frames could have been received at the trunking interface due to the following reasons. The MAC address table of the switches (i) did not contain

**Table 3.4.** Distribution of top applications used in the network

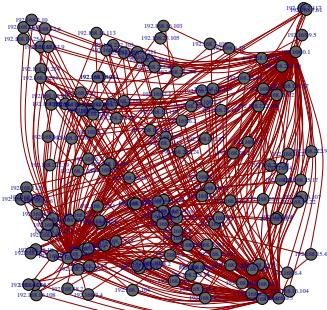
TCP port	Frequency	Application	UDP port	Frequency	Application
13000	11599	Kaspersky	53	27960	DNS
443	10128	HTTPS	137	6589	NetBIOS
88	9700	Kerberos	389	1350	LDAP
80	6938	HTTP	15000	804	Kaspersky NTP
445	5581	Microsoft SMB	88	166	Kerberos
135	940	MS RPC	123	146	NTP
389	880	LDAP	13000	50	Kaspersky
49155	694	Microsoft-DC	138	21	NetBIOS
139	682	Netbios	443	16	QUIC

a record for the destination MAC address and so it was flooded to all ports except the port from which the frame was received; (ii) wrong mapping for the destination MAC address and outgoing interfaces could have existed and so it was forwarded into the wrong port of the switch and thus was received at the trunk port of the router. The two invalid MAC addresses we have discovered are: **00:17:c8:03:a0:7b** and **00:c0:ee:9a:5a:85**. We leave this investigation to the network administrators.

Our next step of data preprocessing was to find for every packet being forwarded a corresponding pair: Note, our trace contained two copies of a packet for which the source and destination IP addresses were within the *192.168.0.0/16* subnetwork with a difference that TTL was reduced by one for one of the packets, and the Ethernet header was recalculated. Thus, by filtering out the duplicates we were able to exclude the possibility of overcounting the number of bytes carried in TCP and UDP streams, and other transport protocols. To perform this filtering step, we merely found all packets whose source and destination addresses were within the subnetwork *192.168.0.0/16* and filtered out frames for which source MAC address was not equal to the MAC address of the VLAN server (note, all VLAN interfaces were assigned the same MAC address). Thus, effectively leaving only one copy of the packet being forwarded between the subnetworks in the trace. Our resulting trace was reduced by 35.4%, and now contained 12190001 frames.

Once again we have searched for frames with invalid MAC addresses (we have introduced the term invalid MAC addresses in the previous paragraphs). In Table 3.3 we list these MAC addresses as well as corresponding IP addresses found in such frames, and manufacturer's name.


Obviously, such discrepancy is strange.



**Figure 3.2.** Interaction between computers in different VLANs over TCP

For the cleaned data, our first step was to take a look at the distribution of TCP and UDP applications and the number of bytes these applications transmitted. Thus, in Table 3.4 we show the distribution of 9 most used UDP and TCP applications. We examined both inter-VLAN and Internet traffic. We observed that while most of hosts were not actively connected to the Internet, using the VPN connections, 15366 connection attempts were made towards the Internet (only TCP SYN packet was captured, which was dropped by the firewall). This could be an indication that quite a lot of connections are made in fact in a stealth mode on users' computers. Next, we have computed the adjacency graph for computers which were interacting over TCP in the local network (an interaction here means that at least one TCP/UDP connection between the pair of computers existed). As it was expected, only few computers (servers) were accepting all TCP connection, with a small fraction of computers interacting between each other. Upon examining closely the logs we have found few interesting things: (i) at least one machine had misconfigured address (it was using self generated IP addresses with the prefix 169.254.171.0/24); (ii) only two machines were interacting with non-server computers 192.168.23.13 and 192.168.3.6 using ports 139 and 445.

Next we computed the distribution of the packet sizes for packets with valid IP addresses. Interestingly there were frames which were larger than maximum allowed Ethernet frames size (so called jumbo frames). Around 10% of the IPv4 packet were larger than 1500 bytes. Finally, we have built the distribution of destination ports used in the packets. In Figure 3.3 we show this distribution.



**Figure 3.3.** Distribution of ports used in the network

### 3.3 Watch out! We are plotting the network map

### 3.4 Performance, performance and once again performance: Stressing the outside world

### 3.5 Looking for the world: CGNs and TCP



## 4. Conclusions

In this short report we have played a bit with the packets which were captured in a small enterprise. Our primary goal was to analyse the interactions of computers and build the statistics for the traffic which was captured for several hours. Our key findings are the following: (i) major traffic in the network is HTTPS and HTTP, (ii) antivirus solutions consume considerable amount of bandwidth, (iii) computers in the network mainly interact with the default gateway and few servers, such as NFS server and mail server.



# **Bibliography**